

Atlas

A personal AI execution engine for food, travel, shopping, and rides
— with approvals before spend, and memory that only personalizes
when you ask it to suggest.

PROBLEM

Chatbots answer. Life still needs twelve apps — and a credit card you trust.

Action is fragmented

Ordering, booking, buying, and scheduling live in separate products with separate carts and payments.

Agents over-personalize

Most memory systems either forget you or push yesterday's favorite every time — even when you asked to explore.

Trust gap on spend

Users won't let an agent spend without clear approvals, constraints, and receipts.

Frontier models are the bottleneck

Great planning, tool use, and recommendation reasoning need frontier-class models — credits unlock iteration speed.

SOLUTION

Execution engine. Chat is just the interface.

- Every request becomes a durable Execution: plan → tools → approve → learn
- Intent-aware memory: conversational / recommendation / execution / hybrid / ambiguous
- Recommendation Engine balances familiarity with exploration — and explains why
- Human-in-the-loop approvals before any payment or booking completes
- MCP tools for live search, food ordering, browse automation, travel, and more

HOW IT WORKS

Intent → domain → the right memory → tools → reply.

Understand

Classify intent

Detect domain

Safety memory

Preference memory

Recommend

Tools

Compose

“Order biryani”

Execution intent. Safety constraints only (allergies, diet, budget). No favorite-pushing.

“What should I eat?”

Recommendation intent. Domain prefs + exploration + live ratings. Why behind each pick.

Ambiguous (“I’m hungry”) → clarify once. One-off orders stay low-confidence and never overwrite long-term prefs.

Differentiation that needs frontier reasoning.

Execution-first

Durable plans, step runners, post-approval resume — not a chat log with tools bolted on.

Trust by design

Approvals, receipts, safety memory (allergies, visa, spending limits) before irreversible actions.

Human-like memory

Confidence-weighted prefs. Suggest only when asked. Explore beyond historical favorites.

These loops are tool-heavy and multi-step. Quality jumps with frontier models for planning, intent classification, recommendation briefings, and tool selection.

STATUS

Working product. Architecture ready to scale.

Shipped

Next.js app: chat, voice, food MCP flow, approvals, activity/receipts, profile, admin routing. Phase 1 execution engine with intent → domain → memory → recommendation pipeline.

What we're validating next

Recommendation quality, tool reliability across domains, approval conversion, and end-to-end food/travel completion rates under frontier models.

Early stage — no vanity metrics. We're optimizing for completed executions, not message volume.

Frontier model access or free API credits.

Not raising cash in this ask — raising capability.

We want sponsored access to frontier models (and/or free API credits) so Atlas can run high-quality multi-step executions, recommendation reasoning, and tool loops without burning the founding budget on tokens.

- **Preferred:** credits or partner access for GPT-class / Claude-class / Gemini-class frontier models
- **Also useful:** embedding + speech credits for memory and voice
- **In return:** case studies, demos, public attribution, feedback on agent/tool-use workloads, and roadmap visibility

CLOSE

Remember like a human. Act with permission.

Help us ship the personal execution layer for everyday life —
powered by frontier models that can plan, choose, and explain.

Next step

Intro call + sandbox access. We'll share architecture, live demo, and a credits
usage plan (eval workloads, monthly ceiling, reporting).